

Nash Bargaining SQP+OptNet Layer: Mathematical Formulation

1 Decision Variables

For a family with M members, each member m has at most T activities. The per-activity decision vector is:

$$\mathbf{x}_{m,t} = [s_{m,t}, e_{m,t}, p_{m,t,1}^{\text{mode}}, \dots, p_{m,t,K}^{\text{mode}}, j_{m,t}, d_{m,t}] \in \mathbb{R}^V, \quad (1)$$

where $V = 2 + K + 2 = 15$ (with $K = 11$ travel modes), and:

- $s_{m,t}, e_{m,t}$: start and end time in z-score normalization ($\mu = 12.946$ h, $\sigma = 4.691$ h),
- $p_{m,t,k}^{\text{mode}}$: probability of travel mode k , forming a simplex $\sum_{k=1}^K p_{m,t,k}^{\text{mode}} = 1$,
- $j_{m,t}$: joint travel probability ($\in [\varepsilon, 1 - \varepsilon]$),
- $d_{m,t}$: driver probability ($\in [\varepsilon, 1 - \varepsilon]$), where $\varepsilon = 10^{-3}$.

The full family decision vector is obtained by concatenating all per-member, per-activity vectors:

$$\mathbf{x} = [\mathbf{x}_{1,1}, \dots, \mathbf{x}_{1,T}, \dots, \mathbf{x}_{M,1}, \dots, \mathbf{x}_{M,T}] \in \mathbb{R}^n, \quad n = M \times T \times V. \quad (2)$$

2 Utility Function

The utility of member m is defined as:

$$\begin{aligned} U_m(\mathbf{x}) = & \beta_0 + \alpha_{\text{social}} \sum_{t=1}^T \tilde{o}_{m,t} a_{m,t} + \alpha_{\text{joint}} \sum_{t=1}^T j_{m,t} a_{m,t} \\ & + \theta_{\text{escort}} \sum_{t=1}^T d_{m,t} (e_{m,t} - s_{m,t}) a_{m,t} \mathbb{I}_{\text{adult}}(m) \\ & + \sum_{t=1}^T \tilde{o}_{m,t} \left(\sum_{k=1}^K p_{m,t,k}^{\text{mode}} \theta_{\text{travel},k} \mathbb{E}[\text{TT}_{m,t,k}] \right) a_{m,t}, \end{aligned} \quad (3)$$

where:

- β_0 : learnable baseline utility,
- α_{social} : learnable social (out-of-home) preference weight,
- α_{joint} : learnable joint travel preference weight,
- θ_{escort} : learnable escort burden coefficient (typically negative),
- $\theta_{\text{travel},k}$: learnable mode-specific travel time disutility (vector of length K),
- $\tilde{o}_{m,t} = 1 - P(\text{dest}_{m,t} = \text{home})$: out-of-home probability from the destination head (fixed, not optimized),

- $\mathbb{E}[\text{TT}_{m,t,k}]$: expected travel time for mode k at activity t , derived from the travel time matrix and destination probabilities (fixed),
- $a_{m,t}$: activity validity mask (1 if activity t of member m is valid, 0 otherwise),
- $\mathbb{K}_{\text{adult}}(m)$: indicator for adult members.

Lower bound. To ensure positivity for the log-utility in the Nash welfare function, a smooth lower bound is applied:

$$U_m \leftarrow U_m + \text{softplus}(1 - U_m), \quad (4)$$

where $\text{softplus}(z) = \log(1 + e^z)$. This guarantees $U_m \geq 1$ smoothly without gradient discontinuity. Utilities of invalid members are set to β_0 to avoid affecting the Nash objective.

3 Nash Welfare Objective

The Nash Bargaining solution maximizes the product of member utilities, equivalently the sum of log-utilities:

$$\max_{\mathbf{x}} W(\mathbf{x}) = \sum_{m=1}^M \log U_m(\mathbf{x}). \quad (5)$$

The objective $W(\mathbf{x})$ is **nonlinear**: the $\log(\cdot)$ composition is inherently nonlinear, and the escort term in U_m contains the bilinear product $d_{m,t}(e_{m,t} - s_{m,t})$.

4 Phase 1: Nonlinear Program

Phase 1 solves the full nonlinear program (NLP) to find a converged point $\mathbf{x}_{\text{conv}}$.

4.1 Optimization Problem

$$\max_{\mathbf{x}} W(\mathbf{x}) = \sum_{m=1}^M \log U_m(\mathbf{x}) \quad (6)$$

subject to the constraints in Table 1.

Table 1: Phase 1 constraints in variable space $\mathbf{x}$.		
Constraint	Formulation	Type
(C1) Variable bounds	$x_i^{\text{lb}} \leq x_i \leq x_i^{\text{ub}}$	Linear
(C2) Time ordering	$e_{m,t} \geq s_{m,t}, \quad s_{m,t+1} \geq e_{m,t}$	Linear
(C3) Mode simplex	$\sum_{k=1}^K p_{m,t,k}^{\text{mode}} = 1$	Linear

The variable bounds (C1) are:

$$s_{m,t}, e_{m,t} \in [z_{\min}, z_{\max}], \quad z_{\min} = \frac{1 - \mu}{\sigma}, \quad z_{\max} = \frac{24 - \mu}{\sigma}, \quad (7)$$

$$p_{m,t,k}^{\text{mode}} \in [\varepsilon, 1 - \varepsilon], \quad (8)$$

$$j_{m,t}, d_{m,t} \in [\varepsilon, 1 - \varepsilon]. \quad (9)$$

Scope of Phase 1 constraints. Phase 1 handles only the simple per-variable and per-activity constraints (C1–C3) via projection. More complex family-level coordination constraints—vehicle capacity and joint travel consistency—involve nonlinear coupling between members and are *not* enforced during Phase 1. These constraints are introduced in Phase 2 (§5), where they are linearized at $\mathbf{x}_{\text{conv}}$ and incorporated as explicit linear constraints in the QP formulation.

4.2 SQP Algorithm

Phase 1 solves the NLP using Sequential Quadratic Programming (SQP) with BFGS inverse Hessian approximation. The constraints (C1–C3) are handled via projection after each step.

Algorithm 1 Phase 1: SQP with BFGS and Armijo Line Search

Require: Decoder output $\boldsymbol{\theta}$, masks, utility parameters

Ensure: Converged point $\mathbf{x}_{\text{conv}}$

```

1:  $\mathbf{x} \leftarrow \boldsymbol{\theta}$                                 ▷ Initialize at decoder output
2:  $H \leftarrow I$                                 ▷ Initial inverse Hessian approximation
3: for  $i = 0, \dots, I_{\text{max}} - 1$  do
4:    $\mathbf{g} \leftarrow -\nabla_{\mathbf{x}} W(\mathbf{x})$                 ▷ Negative gradient of Nash welfare
5:   if  $\|\mathbf{g}\| < \epsilon_{\text{tol}}$  then break                ▷ Converged
6:   end if
7:   if  $i > 0$  then
8:      $\mathbf{s} \leftarrow \mathbf{x}^{(i)} - \mathbf{x}^{(i-1)}, \quad \mathbf{y} \leftarrow \mathbf{g}^{(i)} - \mathbf{g}^{(i-1)}$ 
9:      $H \leftarrow \text{BFGS-UPDATE}(H, \mathbf{s}, \mathbf{y})$         ▷ With Powell's damping
10:  end if
11:   $\mathbf{p} \leftarrow -(H + \delta I) \mathbf{g}$                 ▷ Search direction;  $\delta = 10^{-3}$ 
12:   $\alpha \leftarrow \text{ARMIJO-SEARCH}(\mathbf{x}, \mathbf{p})$         ▷ Backtracking line search
13:   $\mathbf{x} \leftarrow \text{PROJECT}(\mathbf{x} + \alpha \mathbf{p})$         ▷ Project onto (C1)–(C3)
14: end for
15: return  $\mathbf{x}_{\text{conv}} \leftarrow \mathbf{x}$ 

```

BFGS inverse Hessian update with Powell's damping. The matrix H approximates $[\nabla^2(-W)]^{-1}$, updated via the damped BFGS formula. Given step $\mathbf{s} = \mathbf{x}^{(i)} - \mathbf{x}^{(i-1)}$ and gradient change $\mathbf{y} = \mathbf{g}^{(i)} - \mathbf{g}^{(i-1)}$:

1. If $\mathbf{s}^\top \mathbf{y} < 0.2 \mathbf{s}^\top H \mathbf{s}$, apply Powell's damping:

$$\phi = \frac{0.8 \mathbf{s}^\top H \mathbf{s}}{\mathbf{s}^\top H \mathbf{s} - \mathbf{s}^\top \mathbf{y}}, \quad \mathbf{y}_d = \phi \mathbf{y} + (1 - \phi) H \mathbf{s}. \quad (10)$$

Otherwise $\mathbf{y}_d = \mathbf{y}$.

2. Apply the BFGS update:

$$H^+ = \left(I - \rho \mathbf{s} \mathbf{y}_d^\top \right) H \left(I - \rho \mathbf{y}_d \mathbf{s}^\top \right) + \rho \mathbf{s} \mathbf{s}^\top, \quad \rho = \frac{1}{\mathbf{s}^\top \mathbf{y}_d}. \quad (11)$$

Armijo line search. Starting from $\alpha = 1$, backtrack with factor 0.5 until the sufficient decrease condition holds:

$$-W(\text{PROJECT}(\mathbf{x} + \alpha \mathbf{p})) < -W(\mathbf{x}) - c \alpha \|\mathbf{p}\|^2, \quad (12)$$

where $c = 10^{-4}$ is the Armijo parameter.

Projection. The projection operator $\text{PROJECT}(\cdot)$ enforces the linear constraints (C1)–(C3):

1. Clamp times: $s_{m,t}, e_{m,t} \in [z_{\min}, z_{\max}]$.
2. Enforce temporal ordering: for $t = 1, \dots, T$ sequentially, set $s_{m,t} \geq e_{m,t-1}$ and $e_{m,t} \geq s_{m,t}$.
3. Project mode probabilities onto the simplex: clamp to $[\varepsilon, \infty)$ and renormalize.
4. Clamp $j_{m,t}, d_{m,t} \in [\varepsilon, 1 - \varepsilon]$.

5 Phase 2: Quadratic Program

Phase 2 constructs a QP around the converged point $\mathbf{x}_{\text{conv}}$ from Phase 1. The displacement variable is $\mathbf{d} = \mathbf{x} - \mathbf{x}_{\text{conv}}$. The QP enforces two groups of constraints: (i) the Phase 1 constraints (C1–C3) re-expressed in the displacement space, and (ii) **new** family-level coordination constraints (C4–C5) that are introduced for the first time in Phase 2.

5.1 QP Formulation

$$\min_{\mathbf{d}} \frac{1}{2} \mathbf{d}^\top Q \mathbf{d} + \mathbf{p}^\top \mathbf{d} \quad \text{s.t.} \quad G\mathbf{d} \leq \mathbf{h}, \quad A\mathbf{d} = \mathbf{b}. \quad (13)$$

Quadratic term Q . Q is a diagonal matrix encoding the Nash Hessian approximation plus anchor regularization:

$$Q = \text{diag} \left(\frac{1}{U_m^2} + \lambda \right)_i, \quad (14)$$

where U_m is the utility of member m evaluated at $\mathbf{x}_{\text{conv}}$ (Eq. 3), the index i maps to member m via the variable layout, and λ is the fixed anchor weight. The term $1/U_m^2$ approximates the diagonal of the Nash welfare Hessian $\nabla^2 W$, and λ penalizes deviation from the decoder output $\boldsymbol{\theta}$.

Linear term $\mathbf{p}$.

$$\mathbf{p} = \lambda (\mathbf{x}_{\text{conv}} - \boldsymbol{\theta}), \quad (15)$$

where $\boldsymbol{\theta}$ is the decoder output. This term anchors the QP solution toward $\boldsymbol{\theta}$.

Nash correction weight. The relative influence of Nash coordination versus decoder fidelity is:

$$\gamma_{\text{Nash}} = \frac{1/U_m^2}{1/U_m^2 + \lambda}. \quad (16)$$

5.2 Inequality Constraints: $G\mathbf{d} \leq \mathbf{h}$

Inherited from Phase 1 (re-expressed in $\mathbf{d}$)

(C1') Variable bounds. Each variable must stay within its feasible range:

$$-d_i \leq x_i^{\text{conv}} - x_i^{\text{lb}}, \quad d_i \leq x_i^{\text{ub}} - x_i^{\text{conv}}. \quad (17)$$

This yields $2n$ inequality rows.

(C2') Time ordering. The temporal constraints $e_{m,t} \geq s_{m,t}$ and $s_{m,t+1} \geq e_{m,t}$ become:

$$d_{s_{m,t}} - d_{e_{m,t}} \leq e_{m,t}^{\text{conv}} - s_{m,t}^{\text{conv}}, \quad (18)$$

$$d_{e_{m,t}} - d_{s_{m,t+1}} \leq s_{m,t+1}^{\text{conv}} - e_{m,t}^{\text{conv}}. \quad (19)$$

These remain linear, as in Phase 1.

New family-level coordination constraints (introduced in Phase 2)

(C4) Vehicle capacity. This constraint is introduced for the first time in Phase 2. The vehicle capacity constraint $\sum_{m,t} d_{m,t} \cdot p_{m,t}^{\text{car}} \cdot \mathbb{K}_{\text{adult}}(m) \cdot a_{m,t} \leq n_{\text{veh}}$ involves the bilinear product $d_{m,t} \cdot p_{m,t}^{\text{car}}$ (both are decision variables), which cannot be directly included in a QP. It is linearized via first-order Taylor expansion around $\mathbf{x}_{\text{conv}}$:

$$\sum_{m,t} \left(p_{m,t}^{\text{car},0} \Delta d_{m,t} + d_{m,t}^0 \Delta p_{m,t}^{\text{car}} \right) \mathbb{K}_{\text{adult}}(m) a_{m,t} \leq n_{\text{veh}} - \sum_{m,t} d_{m,t}^0 p_{m,t}^{\text{car},0} \mathbb{K}_{\text{adult}}(m) a_{m,t}, \quad (20)$$

where superscript 0 denotes values at $\mathbf{x}_{\text{conv}}$, and Δ denotes the displacement component. This linearization is valid locally around $\mathbf{x}_{\text{conv}}$, which is why this constraint is enforced in Phase 2 (after convergence) rather than Phase 1. It contributes a single linear inequality row.

5.3 Equality Constraints: $\text{Ad} = \mathbf{b}$

Inherited from Phase 1 (re-expressed in $\mathbf{d}$)

(C3') Mode probability normalization. Since $\sum_k p_{m,t,k}^{\text{mode}} = 1$ holds at $\mathbf{x}_{\text{conv}}$, the displacement must preserve this:

$$\sum_{k=1}^K \Delta p_{m,t,k}^{\text{mode}} = 0, \quad \forall m, t. \quad (21)$$

New family-level coordination constraint (introduced in Phase 2)

(C5) Joint travel consistency. This constraint is introduced for the first time in Phase 2. When two members m_1, m_2 are both engaged in joint travel at activity t , their activity attributes should agree. The activation is determined by a sigmoid gate evaluated at $\mathbf{x}_{\text{conv}}$ and fixed as a binary decision:

$$w_{m_1, m_2, t} = \mathbb{K}[\sigma(10(j_{m_1, t}^0 \cdot j_{m_2, t}^0 - \tau)) > 0.1], \quad (22)$$

where $\tau = 0.5$ is the joint consistency threshold, $\sigma(\cdot)$ is the sigmoid function, and superscript 0 denotes values at $\mathbf{x}_{\text{conv}}$. The sigmoid of the bilinear product $j_{m_1, t} \cdot j_{m_2, t}$ is inherently nonlinear; evaluating it at $\mathbf{x}_{\text{conv}}$ and fixing the result as a binary gate converts this into a fixed set of linear equalities suitable for the QP.

For each active pair ($w_{m_1, m_2, t} = 1$):

$$\Delta s_{m_1, t} - \Delta s_{m_2, t} = s_{m_2, t}^0 - s_{m_1, t}^0, \quad (23)$$

$$\Delta e_{m_1, t} - \Delta e_{m_2, t} = e_{m_2, t}^0 - e_{m_1, t}^0, \quad (24)$$

$$\Delta p_{m_1, t, k}^{\text{mode}} - \Delta p_{m_2, t, k}^{\text{mode}} = p_{m_2, t, k}^{\text{mode}, 0} - p_{m_1, t, k}^{\text{mode}, 0}, \quad \forall k, \quad (25)$$

$$\Delta j_{m_1, t} - \Delta j_{m_2, t} = j_{m_2, t}^0 - j_{m_1, t}^0. \quad (26)$$

Each active pair contributes $2 + K + 1 = 14$ equality rows. Note that the right-hand sides drive the paired variables toward agreement: if the pair is active and the values differ at $\mathbf{x}_{\text{conv}}$, the QP will force them to converge.

5.4 Phase 1 $\rightarrow$ Phase 2 Connection

Table 2 summarizes the full constraint set in Phase 2.

Table 2: Phase 2 constraint summary. Constraints C1–C3 are inherited from Phase 1; C4–C5 are new in Phase 2.

Constraint	Phase 2 form	Origin	Linearization method
(C1') Bounds	$G\mathbf{d} \leq \mathbf{h}$	Phase 1	Direct translation
(C2') Time ordering	$G\mathbf{d} \leq \mathbf{h}$	Phase 1	Direct translation
(C3') Mode simplex	$A\mathbf{d} = \mathbf{b}$	Phase 1	Conservation (zero-sum)
(C4) Vehicle capacity	$G\mathbf{d} \leq \mathbf{h}$	New in Phase 2	Taylor expansion at $\mathbf{x}_{\text{conv}}$
(C5) Joint consistency	$A\mathbf{d} = \mathbf{b}$	New in Phase 2	Sigmoid gate fixed at $\mathbf{x}_{\text{conv}}$

Algorithm 2 Phase 2: QP Solve with Gradient Re-attachment

Require: $\mathbf{x}_{\text{conv}}$, $\boldsymbol{\theta}$ (decoder output), masks, utility parameters

Ensure: $\mathbf{x}_{\text{final}}$ (with gradient to decoder and utility parameters)

— **Step 1: Construct QP parameters** —

1: $U_m \leftarrow$ Eq. (3) evaluated at $\mathbf{x}_{\text{conv}}$

2: $Q_{\text{diag}} \leftarrow 1/U_m^2 + \lambda$ (expanded per variable via member mapping) ▷ Eq. (14)

3: $\mathbf{p} \leftarrow \lambda(\mathbf{x}_{\text{conv}} - \boldsymbol{\theta})$ ▷ Eq. (15)

— **Step 2: Build constraints (C1'–C3' inherited)** —

4: $(G, \mathbf{h}) \leftarrow$ Eqs. (17)–(20) ▷ C1', C2'

5: $(A, \mathbf{b}) \leftarrow$ Eqs. (21)–(26) ▷ C3'

— **Step 3: Solve QP** —

6: $\mathbf{d}^* \leftarrow \arg \min_{\mathbf{d}} \frac{1}{2} \mathbf{d}^\top Q \mathbf{d} + \mathbf{p}^\top \mathbf{d}$ s.t. $G\mathbf{d} \leq \mathbf{h}$, $A\mathbf{d} = \mathbf{b}$

— **Step 4: Re-attach gradient** —

7: $\mathbf{d}^* \leftarrow \text{DIAGONALKKTBACKWARD}(\mathbf{p}, Q_{\text{diag}}, \mathbf{d}^*)$ ▷ See §5.7

— **Step 5: Assemble final output** —

8: $\mathbf{x}_{\text{final}} \leftarrow \boldsymbol{\theta} + (\mathbf{x}_{\text{conv}} - \boldsymbol{\theta})_{\text{detach}} + \mathbf{d}^*$ ▷ Eq. (27)

9: $\mathbf{x}_{\text{final}} \leftarrow \text{PROJECT}(\mathbf{x}_{\text{final}})$

10: **return** $\mathbf{x}_{\text{final}}$

5.5 QP Solve

5.6 Final Solution Assembly

The final output is composed as:

$$\mathbf{x}_{\text{final}} = \boldsymbol{\theta} + \underbrace{(\mathbf{x}_{\text{conv}} - \boldsymbol{\theta})}_{\text{no gradient}} + \underbrace{\mathbf{d}^*(\boldsymbol{\theta}, Q)}_{\text{carries gradient}}, \quad (27)$$

where the SQP displacement $(\mathbf{x}_{\text{conv}} - \boldsymbol{\theta})$ is treated as a constant (no gradient flows through Phase 1), and $\mathbf{d}^*$ depends on $\boldsymbol{\theta}$ through $\mathbf{p}$ (Eq. 15) and on utility parameters through Q (Eq. 14).

5.7 Diagonal KKT Backward

For the diagonal QP $\min_{\mathbf{d}} \frac{1}{2} \mathbf{d}^\top \text{diag}(Q_{\text{diag}}) \mathbf{d} + \mathbf{p}^\top \mathbf{d}$, the unconstrained closed-form solution is $\mathbf{d}^* = -\mathbf{p}/Q_{\text{diag}}$. We use this diagonal structure to approximate the gradients of the constrained solution:

$$\frac{\partial d_i^*}{\partial p_i} \approx -\frac{1}{Q_{\text{diag},i}}, \quad \frac{\partial d_i^*}{\partial Q_{\text{diag},i}} \approx -\frac{d_i^*}{Q_{\text{diag},i}}. \quad (28)$$

Given upstream gradient $\bar{\mathbf{d}}$ from the loss function, the gradients propagate as:

$$\bar{\mathbf{p}} = -\bar{\mathbf{d}} / Q_{\text{diag}} \longrightarrow \text{gradient to } \boldsymbol{\theta} \text{ (decoder)}, \quad (29)$$

$$\bar{Q}_{\text{diag}} = -\bar{\mathbf{d}} \odot \mathbf{d}^* / Q_{\text{diag}} \longrightarrow \text{gradient to utility parameters via } 1/U_m^2. \quad (30)$$

6 Straight-Through Log Unpack

After optimization, probability variables (p^{mode}, j, d) must be converted back to log-probabilities for the downstream cross-entropy loss. A naive $\log(p)$ introduces a $1/p$ factor in the gradient, causing instability for small probabilities.

The straight-through log estimator is applied:

$$\tilde{p}_k = p_k + (\log p_k - p_k)_{\text{detach}}, \quad (31)$$

so that the forward value equals $\log p_k$ while the backward gradient is the identity ($\partial \tilde{p}_k / \partial p_k = 1$). For binary heads (joint, driver), the scalar p is expanded to $[1 - p, p]$ before applying Eq. (31).

7 Learnable Parameters

Table 3 summarizes all parameters of the Nash Bargaining layer.

Table 3: Parameters of the Nash Bargaining layer.

Parameter	Dim	Initial Value	Description
β_0	1	2.0	Baseline utility
α_{social}	1	0.3	Out-of-home social preference
α_{joint}	1	0.1	Joint travel preference
θ_{escort}	1	-0.58	Escort burden (adults)
θ_{travel}	$K=11$	$[-1.5, -0.4, \dots]$	Mode-specific travel time disutility
λ (fixed)	1	0.3	Anchor weight (not learned)

All learnable parameters are updated via backpropagation through the diagonal KKT backward (Eqs. 29–30). The anchor weight λ is a fixed hyperparameter controlling the trade-off between Nash coordination and fidelity to the decoder output.

8 Two-Phase Flow Summary

1. **Pack**: Convert decoder predictions (logits) to probability-space vector θ .
2. **Phase 1 (SQP)**: Starting from $\mathbf{x}_0 = \theta$, iterate BFGS + Armijo to find $\mathbf{x}_{\text{conv}}$ (Alg. 1). Handles per-variable constraints (C1–C3) via projection.
3. **Phase 2 (QP)**: Construct QP (Eq. 13) using $\mathbf{x}_{\text{conv}}$ and θ . Enforces all constraints: (C1'–C3') inherited from Phase 1, plus **new** family-level coordination constraints (C4) vehicle capacity and (C5) joint consistency, linearized at $\mathbf{x}_{\text{conv}}$. Gradient flows via diagonal KKT backward (Alg. 2).
4. **Project**: Enforce feasibility on $\mathbf{x}_{\text{final}}$.
5. **Unpack**: Convert back to prediction dict using straight-through log (Eq. 31).

Gradient flows from $\text{loss} \rightarrow \text{unpack} \rightarrow \mathbf{x}_{\text{final}} \rightarrow \mathbf{d}^* \rightarrow$

- $\mathbf{p} \rightarrow \theta \rightarrow$ decoder parameters,
- $Q_{\text{diag}} \rightarrow 1/U_m^2 \rightarrow$ learnable utility parameters.